

WD_Keras-3

December 17, 2020

Create Keras Wide-and-Deep model

This notebook illustrates:

Creating a model using Keras. This requires TensorFlow 2.1

```
[ ]: # Ensure the right version of Tensorflow is installed.  
!pip freeze | grep tensorflow==2.1
```

```
[1]: # change these to try this notebook out  
BUCKET = 'qwiklabs-gcp-02-2449df839737'  
PROJECT = 'qwiklabs-gcp-02-2449df839737'  
REGION = 'us-central1'
```

```
[2]: import os  
os.environ['BUCKET'] = BUCKET  
os.environ['PROJECT'] = PROJECT  
os.environ['REGION'] = REGION
```

```
[3]: %%bash  
if ! gsutil ls | grep -q gs://${BUCKET}/; then  
    gsutil mb -l ${REGION} gs://${BUCKET}  
fi
```

```
[4]: %%bash  
ls *.csv
```

eval.csv
train.csv

0.1 Create Keras model

First, write an input_fn to read the data.

```
[5]: import shutil  
import numpy as np  
import tensorflow as tf  
print(tf.__version__)
```

2.3.1

```
[6]: # Determine CSV, label, and key columns
CSV_COLUMNS = 'weight_pounds,is_male,mother_age,plurality,gestation_weeks,key'.
    ↪split(',')
LABEL_COLUMN = 'weight_pounds'
KEY_COLUMN = 'key'

# Set default values for each CSV column. Treat is_male and plurality as
    ↪strings.
DEFAULTS = [[0.0], ['null'], [0.0], ['null'], [0.0], ['nokey']]
```

```
[7]: def features_and_labels(row_data):
    for unwanted_col in ['key']:
        row_data.pop(unwanted_col)
    label = row_data.pop(LABEL_COLUMN)
    return row_data, label # features, label

# load the training data
def load_dataset(pattern, batch_size=1, mode=tf.estimator.ModeKeys.EVAL):
    dataset = (tf.data.experimental.make_csv_dataset(pattern, batch_size,
    ↪CSV_COLUMNS, DEFAULTS)
        .map(features_and_labels) # features, label
        )
    if mode == tf.estimator.ModeKeys.TRAIN:
        dataset = dataset.shuffle(1000).repeat()
    dataset = dataset.prefetch(1) # take advantage of multi-threading; 1=AUTOTUNE
    return dataset
```

Next, define the feature columns. `mother_age` and `gestation_weeks` should be numeric. The others (`is_male`, `plurality`) should be categorical.

```
[64]: ## Build a Keras wide-and-deep model using its Functional API
def rmse(y_true, y_pred):
    return tf.sqrt(tf.reduce_mean(tf.square(y_pred - y_true)))

# Helper function to handle categorical columns
def categorical_fc(name, values):
    orig = tf.feature_column.categorical_column_with_vocabulary_list(name,
    ↪values)
    wrapped = tf.feature_column.indicator_column(orig)
    return orig, wrapped

def build_wd_model(dnn_hidden_units = [64, 32], nembds = 3):
    # input layer
    deep_inputs = {
        colname : tf.keras.layers.Input(name=colname, shape=(), dtype='float32')
        for colname in ['mother_age', 'gestation_weeks']
    }
```

```

wide_inputs = {
    colname : tf.keras.layers.Input(name=colname, shape=(), dtype='string')
    for colname in ['is_male', 'plurality']
}
inputs = {**wide_inputs, **deep_inputs}

# feature columns from inputs
deep_fc = {
    colname : tf.feature_column.numeric_column(colname)
    for colname in ['mother_age', 'gestation_weeks']
}
wide_fc = {}
is_male, wide_fc['is_male'] = categorical_fc('is_male', ['True', 'False',
↳ 'Unknown'])
plurality, wide_fc['plurality'] = categorical_fc('plurality',
    ['Single(1)', 'Twins(2)', 'Triplets(3)',
    'Quadruplets(4)', 'Quintuplets(5)', 'Multiple(2+)'])

# bucketize the float fields. This makes them wide
age_buckets = tf.feature_column.bucketized_column(deep_fc['mother_age'],
    boundaries=np.
↳ arange(15,45,1).tolist())
wide_fc['age_buckets'] = tf.feature_column.indicator_column(age_buckets)
gestation_buckets = tf.feature_column.
↳ bucketized_column(deep_fc['gestation_weeks'],
    boundaries=np.
↳ arange(17,47,1).tolist())
wide_fc['gestation_buckets'] = tf.feature_column.
↳ indicator_column(gestation_buckets)

# cross all the wide columns. We have to do the crossing before we one-hot
↳ encode
crossed = tf.feature_column.crossed_column(
    [is_male, plurality, age_buckets, gestation_buckets],
↳ hash_bucket_size=20000)
deep_fc['crossed_embeds'] = tf.feature_column.embedding_column(crossed,
↳ nembeds)

# the constructor for DenseFeatures takes a list of numeric columns
# The Functional API in Keras requires that you specify:
↳ LayerConstructor()(inputs)
wide_inputs = tf.keras.layers.DenseFeatures(wide_fc.values(),
↳ name='wide_inputs')(inputs)
deep_inputs = tf.keras.layers.DenseFeatures(deep_fc.values(),
↳ name='deep_inputs')(inputs)

```

```

# hidden layers for the deep side
layers = [int(x) for x in dnn_hidden_units]
deep = deep_inputs
for layerno, numnodes in enumerate(layers):
    deep = tf.keras.layers.Dense(numnodes, activation='relu', name='dnn_{}'.format(layerno+1))(deep)
    deep_out = deep

# linear model for the wide side
wide_out = tf.keras.layers.Dense(10, activation='relu', name='linear')(wide_inputs)

# concatenate the two sides
both = tf.keras.layers.concatenate([deep_out, wide_out], name='both')

# final output is a linear activation because this is regression
output = tf.keras.layers.Dense(1, activation='linear', name='weight')(both)
model = tf.keras.models.Model(inputs, output)
model.compile(optimizer='adam', loss='mse', metrics=[rmse, 'mse'])
return model

print("Here is our Wide-and-Deep architecture so far:\n")
model = build_wd_model()
print(model.summary())

```

Here is our Wide-and-Deep architecture so far:

Model: "functional_21"

Layer (type)	Output Shape	Param #	Connected to
gestation_weeks (InputLayer)	[(None,)]	0	
is_male (InputLayer)	[(None,)]	0	
mother_age (InputLayer)	[(None,)]	0	
plurality (InputLayer)	[(None,)]	0	
deep_inputs (DenseFeatures)	(None, 5)	60000	gestation_weeks[0][0]

```


mother_age[0][0]
is_male[0][0]
plurality[0][0]
-----
dnn_1 (Dense)          (None, 64)          384
deep_inputs[0][0]
-----
wide_inputs (DenseFeatures) (None, 71)          0
gestation_weeks[0][0]
is_male[0][0]
mother_age[0][0]
plurality[0][0]
-----
dnn_2 (Dense)          (None, 32)          2080
dnn_1[0][0]
-----
linear (Dense)          (None, 10)          720
wide_inputs[0][0]
-----
both (Concatenate)      (None, 42)          0
dnn_2[0][0]
linear[0][0]
-----
weight (Dense)          (None, 1)           43
both[0][0]
=====
Total params: 63,227
Trainable params: 63,227
Non-trainable params: 0
-----
None

```

We can visualize the DNN using the Keras plot_model utility.

```
[65]: tf.keras.utils.plot_model(model, 'wd_model.png', show_shapes=False,
    ↪rankdir='LR')
```

[65]:



0.2 Train and evaluate

```
[66]: TRAIN_BATCH_SIZE = 32
NUM_TRAIN_EXAMPLES = 10000 * 5 # training dataset repeats, so it will wrap
    ↳ around
NUM_EVALS = 5 # how many times to evaluate
NUM_EVAL_EXAMPLES = 10000 # enough to get a reasonable sample, but not so much
    ↳ that it slows down

trainds = load_dataset('train*', TRAIN_BATCH_SIZE, tf.estimator.ModeKeys.TRAIN)
evalds = load_dataset('eval*', 1000, tf.estimator.ModeKeys.EVAL).
    ↳ take(NUM_EVAL_EXAMPLES//1000)

steps_per_epoch = NUM_TRAIN_EXAMPLES // (TRAIN_BATCH_SIZE * NUM_EVALS)

history = model.fit(trainds,
                    validation_data=evalds,
                    epochs=NUM_EVALS,
                    steps_per_epoch=steps_per_epoch)
```

Epoch 1/5

312/312 [=====] - 3s 10ms/step - loss: 1.4212 - rmse: 1.1644 - mse: 1.4212 - val_loss: 1.2406 - val_rmse: 1.1135 - val_mse: 1.2406

Epoch 2/5

312/312 [=====] - 3s 10ms/step - loss: 1.1908 - rmse: 1.0786 - mse: 1.1908 - val_loss: 1.3444 - val_rmse: 1.1588 - val_mse: 1.3444

Epoch 3/5

312/312 [=====] - 3s 10ms/step - loss: 1.1961 - rmse: 1.0802 - mse: 1.1961 - val_loss: 1.1468 - val_rmse: 1.0706 - val_mse: 1.1468

Epoch 4/5

312/312 [=====] - 2s 7ms/step - loss: 1.1317 - rmse: 1.0519 - mse: 1.1317 - val_loss: 1.1900 - val_rmse: 1.0906 - val_mse: 1.1900


Epoch 5/5

312/312 [=====] - 3s 10ms/step - loss: 1.1233 - rmse: 1.0490 - mse: 1.1233 - val_loss: 1.1407 - val_rmse: 1.0675 - val_mse: 1.1407

0.3 Visualize loss curve

```
[67]: # plot
import matplotlib.pyplot as plt
nrows = 1
ncols = 2
fig = plt.figure(figsize=(10, 5))

for idx, key in enumerate(['loss', 'rmse']):
    ax = fig.add_subplot(nrows, ncols, idx+1)
    plt.plot(history.history[key])
    plt.plot(history.history['val_{}'.format(key)])
    plt.title('model {}'.format(key))
    plt.ylabel(key)
    plt.xlabel('epoch')
    plt.legend(['train', 'validation'], loc='upper left');
```



0.4 Save the model

```
[68]: import shutil, os, datetime
OUTPUT_DIR = 'babyweight_trained'
shutil.rmtree(OUTPUT_DIR, ignore_errors=True)
EXPORT_PATH = os.path.join(OUTPUT_DIR, datetime.datetime.now().
    ↳strftime('%Y%m%d%H%M%S'))
tf.saved_model.save(model, EXPORT_PATH) # with default serving function
print("Exported trained model to {}".format(EXPORT_PATH))
```

INFO:tensorflow:Assets written to: babyweight_trained/20201217213448/assets

Exported trained model to babyweight_trained/20201217213448

```
[69]: !ls $EXPORT_PATH
```

```
assets  saved_model.pb  variables
```

```
Copyright 2020 Google Inc. Licensed under the Apache License,  
Version 2.0 (the ``License''); you may not use this file except in  
compliance with the License. You may obtain a copy of the License at  
http://www.apache.org/licenses/LICENSE-2.0 Unless required by applicable law or  
agreed to in writing, software distributed under the License is distributed on  
an ``AS IS'' BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express  
or implied. See the License for the specific language governing permissions and  
limitations under the License
```